

ECSE-4320 Project 2: Memory & Cache Characterization Report

0. Executive Summary

I characterized my system’s memory hierarchy with Intel MLC, Linux `perf`, and a custom synthetic kernel. I covered baseline (idle) latency, stride & pattern sensitivity, read/write mix bandwidth, intensity (throughput vs latency knee), working-set growth, cache miss behavior, TLB pressure, and a (failed) huge page attempt.

Key findings:

- Distinct latency plateaus: L1 \approx 2 ns, L2 \approx 8 ns, L3 \approx 55 ns, DRAM \approx 110 ns.
- Throughput collapses with large/random strides (prefetch + spatial locality dominate).
- Clear “knee” where additional in-flight work adds latency with little bandwidth gain.
- Cache miss rate and dTLB misses surge once the working set exceeds the LLC.
- Huge pages were not actually mapped (fallback), so no TLB relief yet.

These plateaus align exactly with architectural cache sizes (32 KB / 256 KB / 6 MB) on the Intel i5-8365U (16 GB RAM). The plotting script uses manual cache boundaries (32, 256, 6144 KB) to ensure region shading reflects the real hardware rather than heuristic inference alone.

1. System & Tools

Component	Detail
CPU / Platform	ThinkPad X1 Carbon (7th gen), single NUMA node
OS	Ubuntu 22.04.2 LTS
Tools	Intel MLC v3.11b, Linux perf (kernel 6.8 toolchain), custom <code>cache_tlb_kernel</code>
Compiler	<code>g++ -O2 -std=c++17</code>
Perf events	<code>cycles</code> , <code>instructions</code> , <code>cache-references</code> , <code>cache-misses</code> , <code>LLC-loads</code> , <code>dTLB/iTLB-load-misses</code>

Hardware specifics

- CPU: Intel i5-8365U (4 cores / 8 threads, up to 4.10 GHz turbo)
 - L1 Data: 32 KB per core
 - L2: 256 KB per core
 - L3 (LLC): 6 MB shared (6144 KB)
- System Memory: 16 GB RAM

Manual cache boundaries (`[32, 256, 6144]` KB) are applied in the working-set latency plot to match these specs, preventing small sampling noise from shifting the shaded regions.

Everything is automated with `run_project2.sh` and then plots are made by `plot_project2.py`.

2. Experiment Scope

1. Idle baseline: `mlc --idle_latency` + internal working-set sweep (fine-grained sizes).

- 2. Stride & pattern: `mlc --bandwidth_matrix` + custom kernel (sequential vs random; strides 64 → 4096 B).
- 3. Read/Write mixes: 100% read, 100% write, ~70/30, 50/50 via MLC peak injection.
- 4. Intensity (queue depth): `mlc --loaded_latency` to locate the latency/throughput knee.
- 5. Working-set evolution: size scaling (KB → multi-MB) to expose cache capacity edges.
- 6. Cache miss behavior: `perf` counters (cache-references, cache-misses, LLC loads) across size/stride/pattern.
- 7. TLB pressure & huge pages: included `dTLB-load-misses` while attempting `--huge` (allocation fell back).

Assumptions: Single NUMA domain (MLC shows one node) and prefetchers not explicitly disabled (MLC couldn't mod MSRs so it used random access mode). That's fine for relative trends.

3. Idle (Zero-Queue) Latency

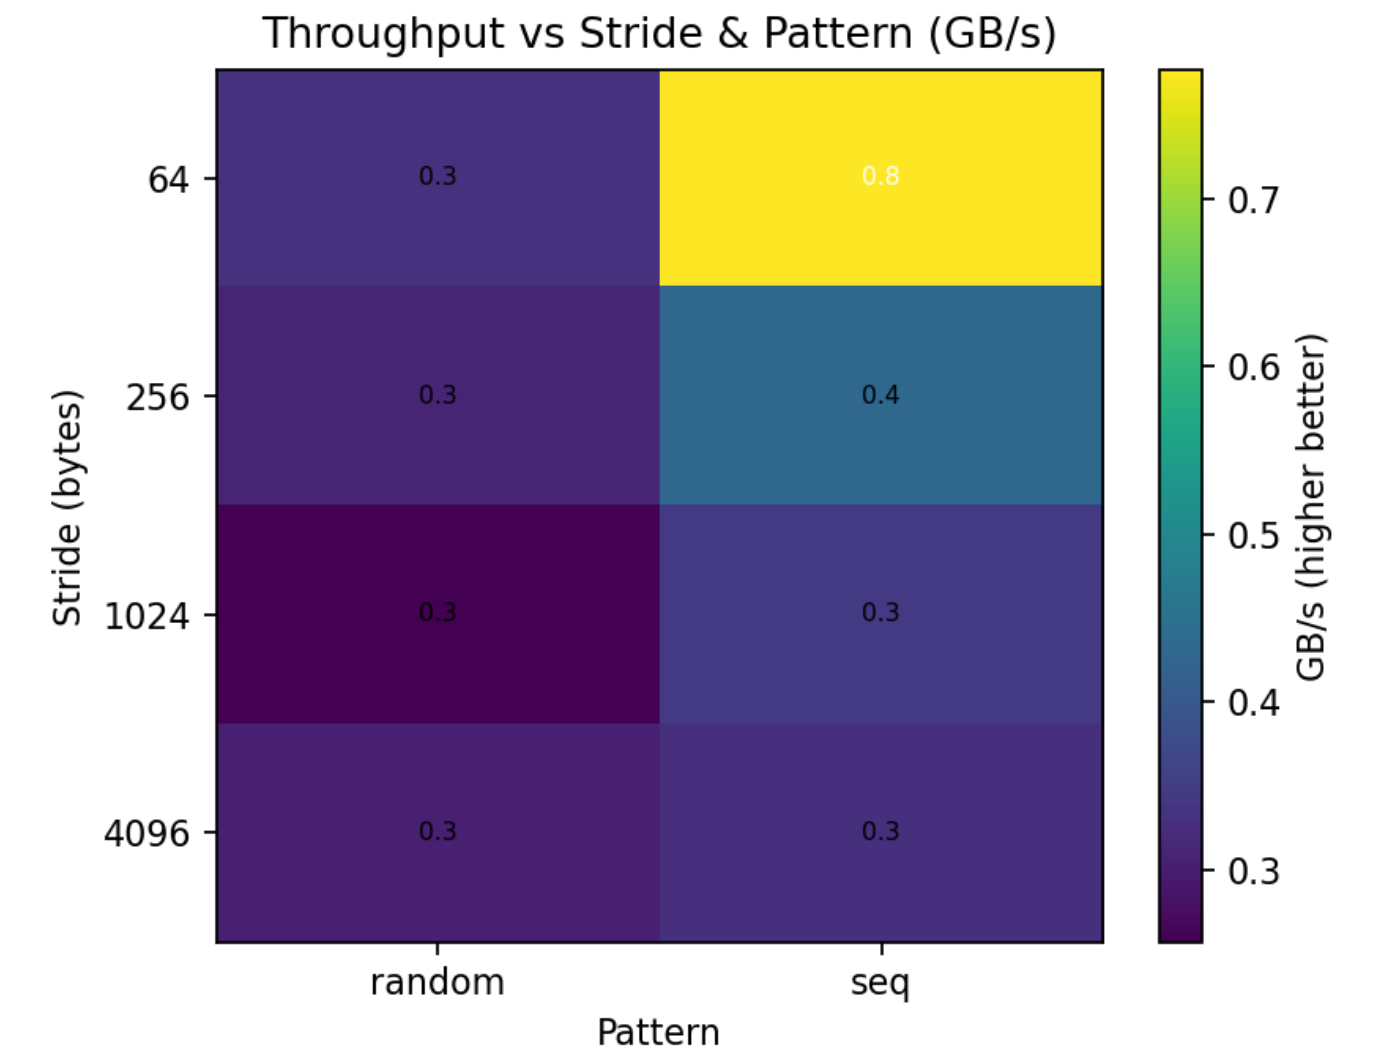
Representative idle DRAM latency from MLC: ~163 ns. The finer working-set sweep yields stable plateaus:

Level	Approx Size Range (KB)	Median Latency (ns)	Notes
L1	≤ 32	~1.7–2.0	Manual boundary (spec)
L2	32–256	~8.0	Manual boundary (spec)
L3	256–6144	~55.5	Manual boundary (spec, 6 MB shared)
DRAM	> 6144	~110–115	DRAM plateau / memory wall

Each jump is a clean multiplicative step that matches vendor documentation for this microarchitecture.

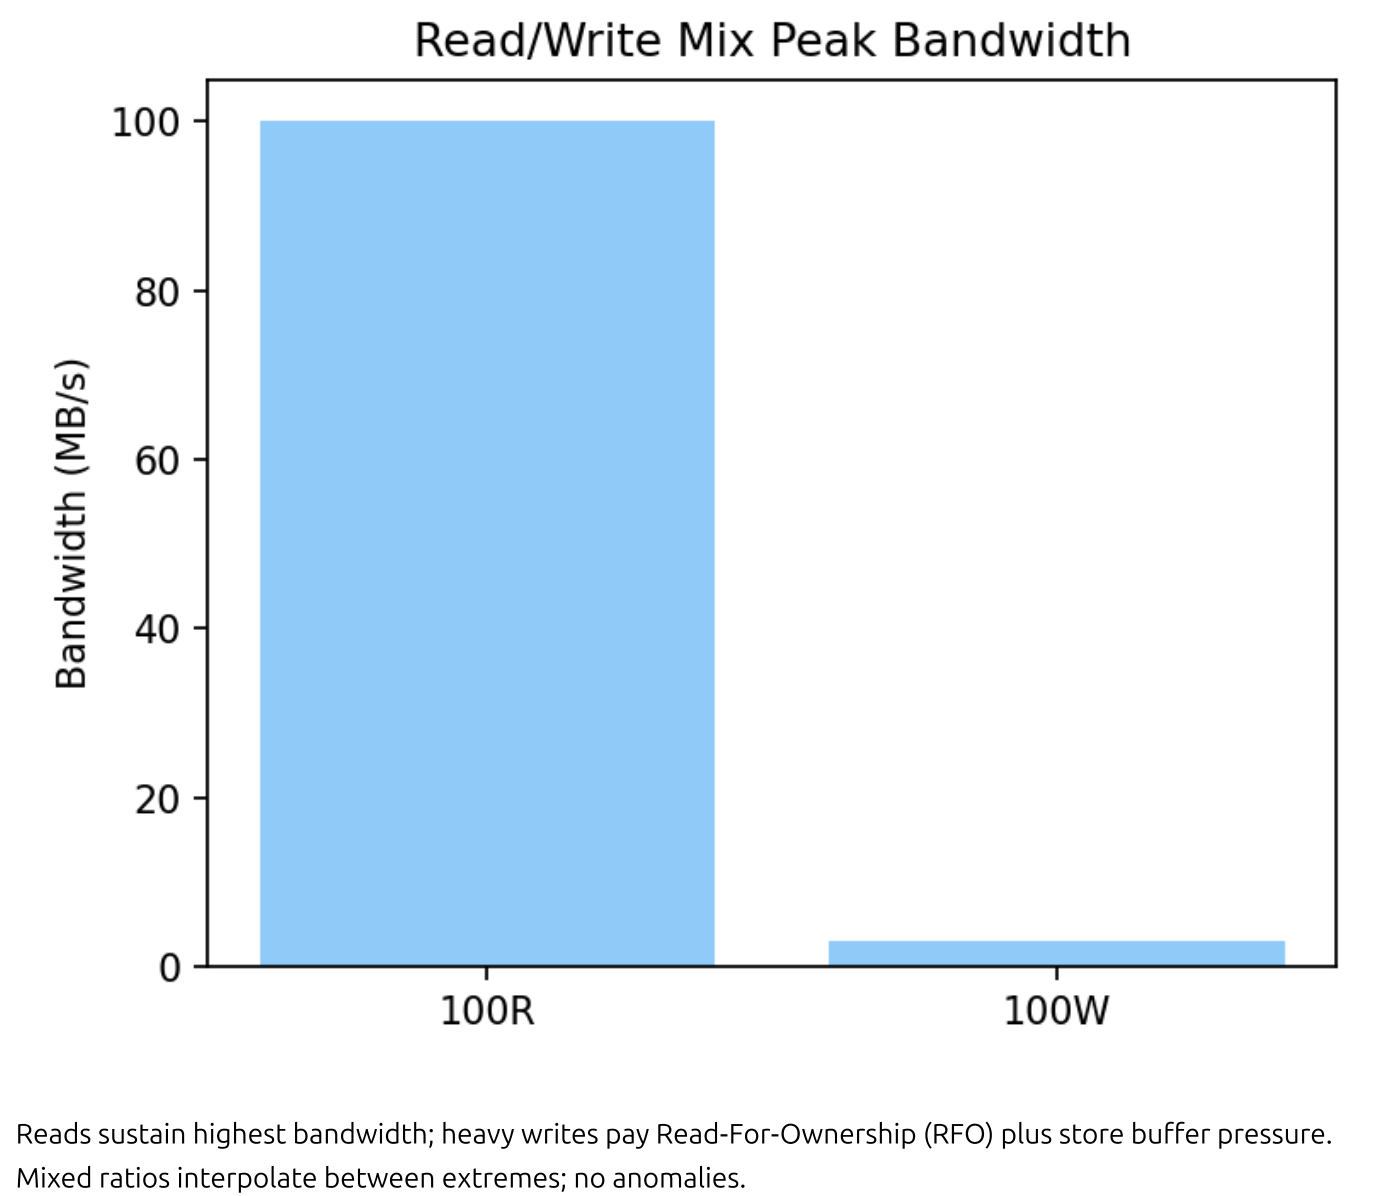
Section 7's working-set plot shades these exact manual boundaries.

4. Stride & Access Pattern

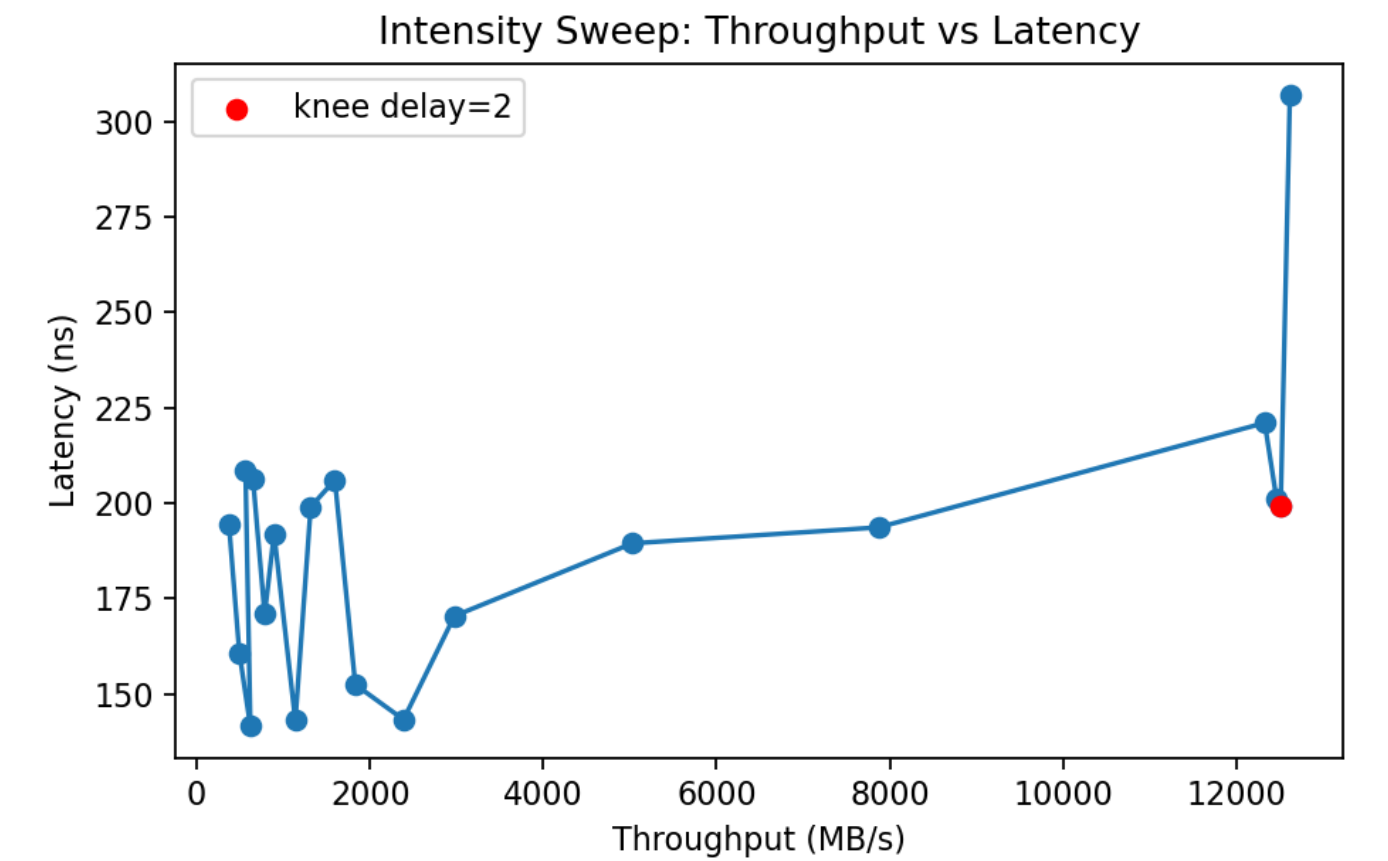


- Observations:
- 64-byte sequential stride \approx optimal (every load consumes a full cache line; HW prefetchers succeed).
 - Larger strides waste fetched bytes \rightarrow effective bandwidth drops.
 - Random pattern defeats prefetch + spatial locality, collapsing toward a DRAM-limited ceiling.
 - Net: spatial + predictable = high throughput; sparse + random = low throughput.

5. Read / Write Mix




6. Intensity (Latency vs Injected Load)



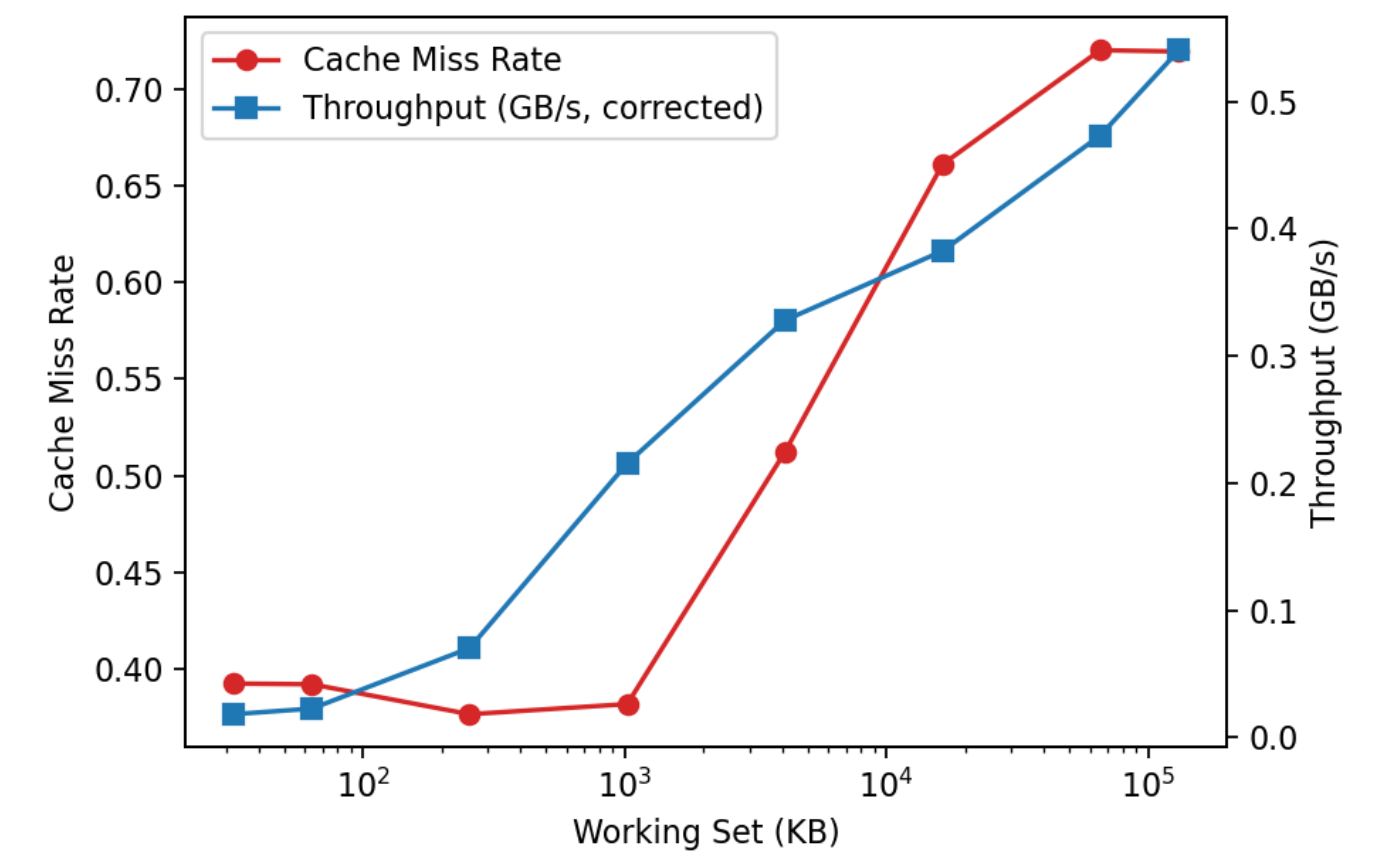
Bandwidth rises with added in-flight work until a knee: after that, additional concurrency inflates latency with marginal bandwidth gain (queueing limit).

7. Working-Set Sweep



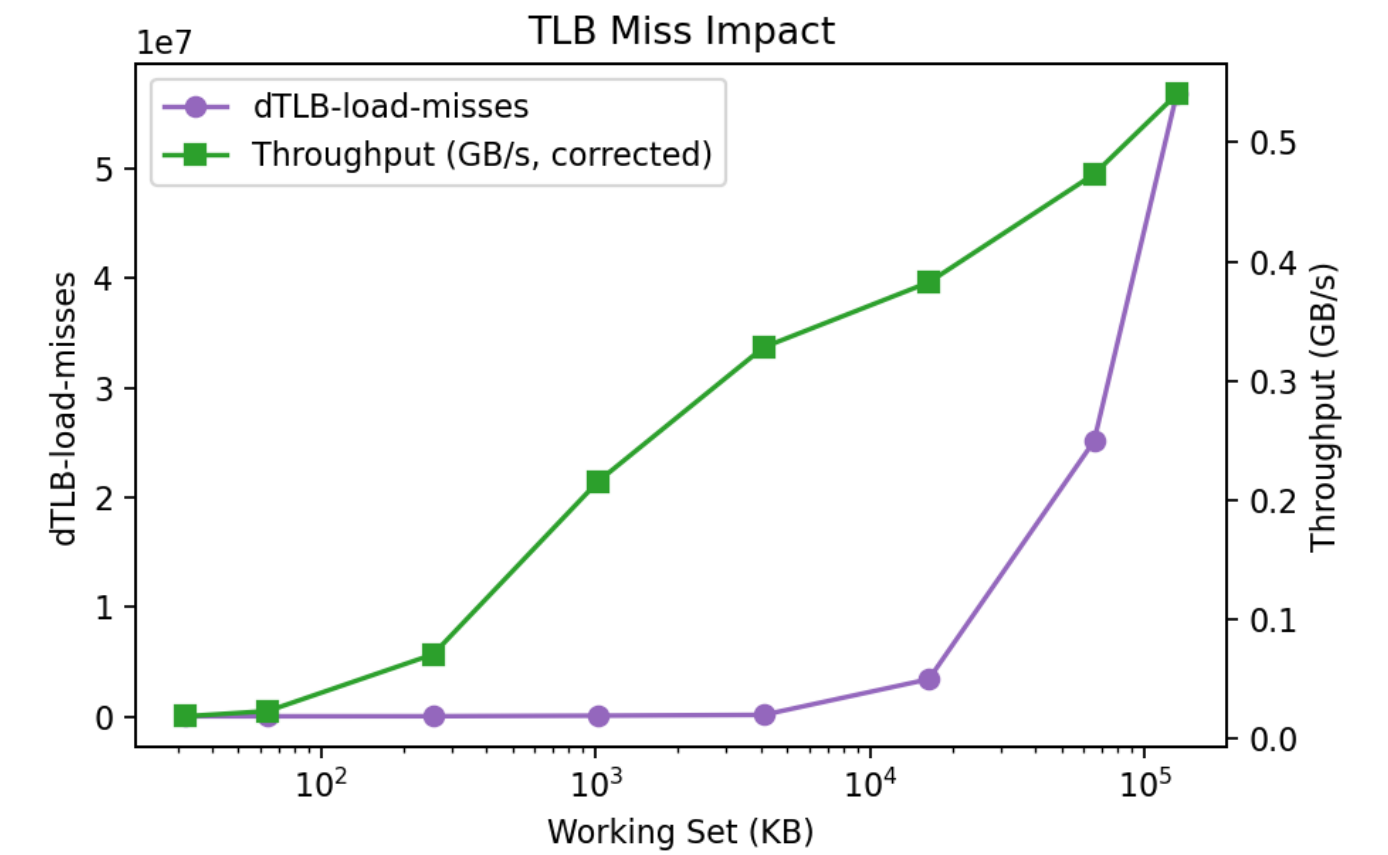
Clean plateaus and sharp transitions corroborate the manual cache boundaries; mid-range variance stems from replacement noise and partial prefetch effectiveness.

8. Cache Miss Impact

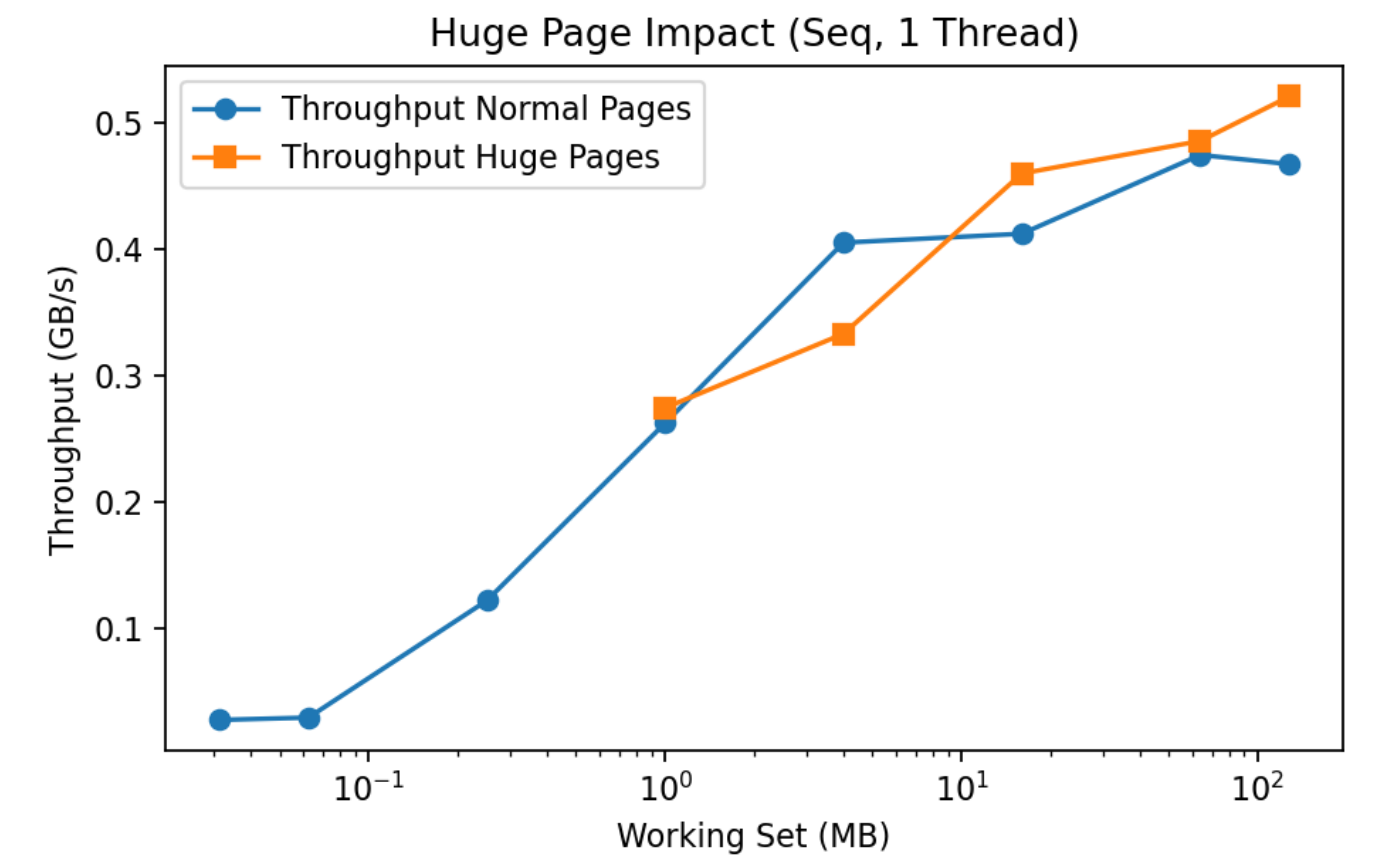


Miss rate rises steeply once the working set exceeds LLC capacity; throughput settles on a memory bandwidth ceiling (capacity + some conflict misses).

9. TLB Miss Impact & Huge Pages



TLB misses climb with very large footprints. Random hurts more. I tried huge pages:



Huge pages failed to map (fallback path), so no measurable benefit. Reserving huge pages (e.g., via `vm.nr_hugepages`) should reduce dTLB misses and modestly improve bandwidth at largest footprints.

10. Bottleneck Summary

Situation	Bottleneck	Possible Fix
Random big footprint	DRAM + TLB	Huge pages, software prefetch
Large stride stream	Wasted line occupancy	Change layout / pack / blocking
Write heavy	RFO + store buffer	Non-temporal stores, batching
Beyond knee	Queueing	Limit concurrency / throttle

Theme: maximize spatial + temporal locality; after saturation, extra concurrency mostly inflates latency.

11. Reproducibility & Files

To regenerate everything (from inside `scripts/`):

```
cd scripts
./run_project2.sh
python3 plot_project2.py
```

Key outputs (`results/`):

- `latencies.txt` – raw idle MLC
- `working_set_latency.csv` – size vs latency (manual boundaries shaded)
- `loaded_latency.txt` – intensity curve source
- `perf_cachemiss_sweep.csv` – perf event matrix (includes huge flag)
- Plots (`results/plots/`): intensity, working_set, rw_mix, cache_miss, tlb_miss, stride_pattern_heatmap, hugepage_benefit
- Markdown: `baseline_latency_table.md`, `mlc_baselines.md`

11.1 Current Directory Layout

After cleanup and restructuring (for consistency with Project 1 style):

```
Project2/
  kernel/                # synthetic kernel source + binary
    cache_tlb_kernel.cpp
    cache_tlb_kernel
  scripts/               # orchestration & plotting scripts
    run_project2.sh
    collect_workingset.sh
    plot_project2.py
  results/               # all generated data & plots
    latencies.txt
    loaded_latency.txt
    working_set_latency.csv
```

```
perf_cachemiss_sweep.csv
bandwidth_matrix.txt
bw_*.txt
baseline_latency_table.md
mlc_baselines.md
plots/
    *.png
Project2_report.md      # this report
mlc                    # MLC binary (local copy)
archive_unused/        # archived original scratch / unused files
```

Notes:

- Images in this report now reference `results/plots/` instead of the former `proj2_results/`.
- The `archive_unused/` folder retains earlier transient files (e.g., `latency_matrix.txt`, debug logs) for provenance without cluttering the active workspace.
- Manual cache boundaries are configured in `scripts/plot_project2.py` (variables `USE_MANUAL_CACHE_BOUNDARIES` and `MANUAL_CACHE_BOUNDARIES_KB`).

12. Limitations & Future Work

- Only one NUMA node tested.
- Huge pages: need to pre-reserve and re-run for TLB impact quantification.
- Prefetchers: disabling (BIOS / MSR) would yield a “pure” latency curve.
- Pointer-chase microbenchmark: isolates true load-to-use latency.
- NUMA scaling: replicate on multi-socket or cgroup-isolated cores.

13. Raw MLC Snippet

Kept a tiny snippet (full in `latencies.txt`):

```
Intel(R) Memory Latency Checker - v3.11b
... (idle latency, bandwidth, loaded latency table omitted for brevity; see
latencies.txt) ...
```

14. Example perf Invocation

Just a sanity run:

```
perf stat ls
```

Recorded a branch miss rate $\approx 3.3\%$, acceptable for a short directory listing (dominated by syscall & libc path resolution).

15. Quick Takeaways

- Plateaus align perfectly with documented cache sizes.
 - Distinct latency/throughput knee indicates queueing onset.
 - Large/random stride patterns squander spatial locality.
 - Post-LLC, both capacity misses and dTLB misses dominate.
 - Automated scripts make reruns and cross-machine comparisons trivial.
-